

# Step-by-Step Implementation Guide

## Intelligence Tokens (ITs) + Object Store / Vector Store Prototype (Template for ASCE Leadership B)

Build date: 2025-12-24

Prepared for: Eva Lerner-Lam (internal prototyping; not official ASCE policy)

### 1. What you are building

You are building a prototype workflow that packages domain recommendations into an Intelligence Token (IT) that:

- (1) references authoritative, immutable artifacts in an Object Store, and
- (2) makes those artifacts discoverable via a Vector Store for retrieval-augmented responses,

while enforcing governance policy gates and preserving attribution metadata suitable for compensation routing.

Key idea:

- The Object Store is the source of truth (canonical artifacts + hashes).
- The Vector Store is the derived index (embeddings + metadata + references back to the Object Store).

### 2. Minimum architecture (conceptual)

Components

- A) Authoring workstation + repository
- B) Object Store (prototype: GitHub repository with immutable history)
- C) Packetization layer (canonicalization + hashing + manifests)
- D) Vector Store (any vendor that supports metadata filtering)
- E) Runtime policy gates (filters + human review hooks)
- F) Audit log / usage ledger (for attribution and compensation evidence)

Data flow

Objects → Manifest → IT object → Chunk + Embed → Vector records → Retrieval with policy gates → Response with attribution

### 3. Step-by-step build plan

Step 1 — Stand up the Object Store (prototype)

- Create a GitHub repository that will host /object\_store and /schemas.
- Use the folder structure provided in 03\_object\_store\_template/.

Step 2 — Define canonicalization and hashing

- Choose one canonicalization method (e.g., sorted keys + minified JSON).
- Compute SHA-256 hashes of canonical content bytes.
- Treat the hash output as the content-addressed identifier (CID) for prototypes.

Step 3 — Create an Object Manifest

- For each object you want searchable, list:  
object\_id, title, type, provenance, content\_hash, and repository path.
- Save it using 02\_schemas/object\_manifest.schema.json.

#### Step 4 — Create the Intelligence Token (IT) object

- Set token\_type = transferable\_knowledge for the recommendations package.
- Reference one or more manifest IDs.
- Specify governance tier and required attribution fields.
- Define policy gates (access/licensing/quality).

#### Step 5 — Convert objects into indexable text

- Extract canonical text from each object (retain originals in Object Store).
- Preserve headings/section labels to maintain governance context.

#### Step 6 — Chunk and embed

- Chunk at paragraph/section level; keep overlap.
- Generate embeddings using one stable embedding model.
- Produce vector records per 02\_schemas/vector\_record.schema.json.

#### Step 7 — Ingest into your Vector Store

- Upsert vector records; ensure metadata filtering is enabled.
- Store policy gate references in each record's metadata.

#### Step 8 — Implement retrieval with policy gates

- Apply gates before similarity search when possible (pre-filter).
- Otherwise filter results after retrieval, before responding.
- Always return attribution fields with any retrieved content.

#### Step 9 — Verification & Validation (V&V)

- Run schema validation and hash verification in CI (continuous integration).
- Confirm provenance completeness.
- Test retrieval filters and audit logging.

#### Step 10 — Demonstrate to ASCE stakeholders

- Use a small set of objects representing governance, ethics, policy, and technical guidance.
- Show: traceability (prompt → retrieved chunks → object IDs → hashes → provenance).

### **4. Verification & Validation checklist (minimum)**

#### Critical checks

- Schema validation passes for all manifests, IT objects, and vector records.
- Hash verification: recomputed SHA-256 matches recorded hashes.
- Provenance: each object contains provenance, author/org, publication date (if known), and license/rights.
- Policy gates: retrieval enforces access and licensing constraints.
- Attribution: responses include citation-ready attribution metadata.
- Audit logging: record prompt context, retrieved record IDs, and policy gate decisions (without storing sensitive data).

### **5. Compensation and attribution (prototype-ready, policy-neutral)**

This prototype does not prescribe ASCE policy. It supplies metadata and logging hooks that could enable compensation programs.

#### Recommended evidence artifacts

- Usage log entry: token\_id, object\_id, chunk\_id, timestamp, user role, gate decision.
- Attribution bundle: author, organization, publication date, license, and canonical object hash.

#### Human review requirement

- Any compensation determination should be gated as “require\_human\_review” until an approved program exists.